

Research Interests	I design human-centric foundation models. My work explores post-training with synthetic data and reward functions from user-modeling simulations [P1, R5, R2], personalized alignment methods [R3, R4, R11], and real-world evaluations with specialized user interfaces [R6, R8].	
Education	University of Southern California, Information Sciences Institute (USC ISI) Ph.D. in Computer Science, Viterbi School of Engineering Advisor: Jonathan May	09/2020 - Present
	Hong Kong University of Science and Technology (HKUST) B.Eng. in Computer Science, School of Engineering	09/2013 - 06/2019
Experience	Amazon AGI , Research Intern Developed Trial-Error-Explain ICL, a tuning-free personalized text generation method [R3].	05/2024 - 09/2024
	Amazon Alexa , Research Intern Post-trained LLMs with RL from human feedback on audio responses for voice-suitable response generation [R5].	06/2023 - 09/2023
	Meta AI , Research Intern Implemented example-guided question answering for continual learning for dialogue state tracking (DST) for on-device models [R9].	06/2022 - 09/2022
	Meta AI , Research Intern Developed CheckDST for comprehensive DST evaluation and PrefineDST for pre-finetuning DST models to improve their robustness to adversarial perturbations [R12].	08/2021 - 03/2022
	Stitch Fix Inc. , Data Science Intern Developed a client feedback and request notes analysis system that identifies key seasonal demand trends. [Blog post]	06/2020 - 08/2020
	USC ISI , Programmer Analyst Post-trained dialogue models for grounded response generation using improvised dialogues [R17].	09/2019 - 06/2020
Publications	*: equal contribution. Featured work	
	[R1] Can Vision Language Models Understand Mimed Actions? <u>Hyundong Cho</u> , S. Lin, T. Srinivasan, M. Saxon, D. Kwon, N. Chavez, J. May. 2025 <i>ACL 2025 Findings</i>	
	[R2] NewsInterview: Evaluating LLMs' Grounding Gap via Informational Interviews M. Lu, S. Kalyan, <u>Hyundong Cho</u> , W. Shi, J. May, A. Spangher. 2025 <i>ACL 2025</i>	
	[R3] Tuning-Free Personalized Alignment via Trial-Error-Explain In-Context Learning <u>Hyundong Cho</u> , K. Sharma, N. Jedema, L. Ribeiro, A. Moschitti, J. May. <i>NAACL 2025 Findings</i>	
	[R4] Aligning Language Models with Demonstrated Feedback O. Shaikh, M. Lam, J. Hejna, Y. Shao, <u>Hyundong Cho</u> , M. Bernstein, D. Yang. 2024 <i>ICLR 2025</i>	

- [R5] [Speechworthy Instruction-tuned Language Models](#)
Hyundong Cho, N. Jedema, L. Ribeiro, K. Sharma, P. Szekely, A. Moschitti, R. Janssen, J. May
EMNLP 2024
- [R6] [Can Language Model Moderators Improve the Health of Online Discourse?](#)
Hyundong Cho, S. Liu, T. Shi, D. Jain, B. Rizk, Y. Huang, Z. Lu, N. Wen, J. Gratch, E. Ferrara, J. May
NAACL 2024
- [R7] [Improving Covert Toxicity Detection by Retrieving and Generating References](#)
D. Lee, Hyundong Cho, W. Jin, J. Moon, S. Park, P. Röttger, J. Pujara, R. Lee
NAACL 2024 Workshop on Online Abuse and Harms
- [R8] [BotEval: Facilitating Interactive Human Evaluation](#)
Hyundong Cho, Thamme Gowda, Y. Huang, Z. Lu, Tianli Tong, J. May
ACL 2024 System Demonstration
- [R9] [Continual Dialogue State Tracking via Example-Guided Question Answering](#)
Hyundong Cho, A. Madotto, Z. Lin, K. Raghavi Chandu, S. Kottur, J. Xu, J. May, C. Sankar
EMNLP 2023
- [R10] [Analyzing Norm Violations in Live-Stream Chat](#)
J. Moon, D. Lee, Hyundong Cho, W. Jin, C. Park, M. Kim, J. May, Jay Pujara, S. Park
EMNLP 2023
- [R11] [RECAP: Retrieval-Enhanced Context-Aware Prefix Encoder for Personalized Dialogue Response Generation](#)
S. Liu, Hyundong Cho, M. Freedman, X. Ma, J. May
ACL 2023
- [R12] [Know Thy Strengths: Comprehensive Dialogue State Tracking Diagnostics](#)
Hyundong Cho, C. Sankar, C. Lin, A. Celikyilmaz, S. Shayandeh, J. May, A. Beirami
EMNLP 2022 Findings
- [R13] [Reflect, Not Reflex: Inference-Based Common Ground Improves Dialogue Response Quality](#)
P. Zhou, Hyundong Cho, P. Jandaghi, D. Lee, Y. Lin, J. Pujara, X. Ren
EMNLP 2022
- [R14] [Probing Commonsense Explanation in Dialogue Response Generation](#)
P. Zhou, P. Jandaghi, Hyundong Cho, Y. Lin, J. Pujara, X. Ren
EMNLP 2021 Findings
- [R15] [Viola: A Topic Agnostic Generate-and-Rank Dialogue System](#)
Hyundong Cho, B. Shbita, K. Shenoy, S. Liu, N. Patel, H. Pindikanti, J. Lee, J. May 2021
Alexa Prize Socialbot Grand Challenge 4 Proceedings
- [R16] [Agenda Pushing in Email to Thwart Phishing](#)
Hyundong Cho, G. Bartlett, M. Freedman
ACL 2021 Doc2Dial Workshop
- [R17] [Grounding Conversations with Improvised Dialogues](#)
Hyundong Cho, J. May
ACL 2020

Preprints

- [P1] [Which Questions Improve Learning the Most? Utility Estimation of Questions with LM-based Simulations](#)
D. Lee*, Hyundong Cho*, J. May, J. Pujara. 2025
- [P2] [Uncovering Intervention Opportunities for Suicide Prevention with Language Model Assistants](#)
J. Ranjit, Hyundong Cho, C. Smerdon, Y. Nam, M. Phung, J. May, J. Blosnich, S. Swayamdipta. 2025

Awards	Alexa Prize Socialbot Grand Challenge 4 Semi-finalist , \$250,000 grant, Amazon, 2021. Graduate School Fellowship for Incoming Students , USC, 2020. First Class Honors , HKUST, 2019. Academic Excellence Recruitment Scholarship , Full tuition, HKUST, 2013-2019. Reaching Out Award , HKSAR Government Scholarship Fund, 2019. CodeIT Suisse Coding Challenge Winner , Credit Suisse APAC Region, 2017/9. Dean's List , HKUST, 2013, 2015, 2017, 2018, 2019.
Invited Talks	CSCI 499 Language Models in NLP : Guest Lecture on RLHF CSCI 662 Advanced NLP : Guest Lecture on RLHF USC ISI NLG Seminar : Contextualization for Human-AI Interactions Tonal Talk : LLMs and RLHF
Teaching Experience	TA , CSCI 662 Advanced NLP TA , CSCI 699 Probabilistic and Generative Models
Professional Services	Reviewer : *ACL, EMNLP, ACL ARR (2021-2025), NeurIPS (2023-2025), ICLR (2025), TMLR (2023-2024) Lead Workshop Organizer : What's left to TEACH chatbots Workshop at ICML 2023